

Machine Learning 1


Pierce Ford (PID: A59010464)

10/22/2021

#Clustering Methods

Kmeans clustering in R is done with the `kmeans()` function. Here we make up some data to test and learn with.

```
tmp <- c(rnorm(30, 3), rnorm(30, -3))  
#Make a two column dataset that includes tmp and tmp reversed, to make two clear  
#groups of points  
data <- cbind(tmp, rev(tmp))  
plot(data)
```



Run `kmeans()` set `k` (centers) to 2 and `nstart` to 20. The thing with Kmeans is you have to tell it how many clusters you want.

```
km <- kmeans(data, centers=2, nstart=20)
km
```

[illegible]

Q. How many points are in each cluster?

km\$size

```
## [1] 30 30
```

Q. What 'component of your result object details cluster assignment/membership?

```
km$cluster
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2  
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Q. What 'component of your result object details cluster cluster center?

km\$centers

```
##          tmp
## 1  3.245010 -3.190075
## 2 -3.190075  3.245010
```

Q. Plot `x` colored by the `kmeans` cluster assignment and add cluster centers as blue points?

```
plot(data, col=km$cluster)
points(km$centers, col="blue", pch=15)
```



#Hierarchical Clustering

We will use the `hclust()` function on the same data as before and see how this method works.


```
hc <- hclust(dist(data))
hc
```

```
##
## Call:
## hclust(d = dist(data))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

Hclust has a plot method

```
plot(hc, cex=0.7)
```

Cluster Dendrogram



```
dist(data)
hclust (*, "complete")
```


To find our membership vector we need to “cut” the tree and for this we use the `cutree()` function and tell it the height to cut at.

```
cutree(hc, h=7)
```

[illegible]

We can also use `cutree()` and state the number of k cluster we want.

```
grps <- cutree(hc, k=2)
#plot with hc clusters as color
plot(data, col=grps)
```



#Principal Component Analysis (PCA)

PCA is useful for visualizing key variance in datasets with high dimensionality.

##PCA of UK Food Data Let's read in the UK food dataset.


```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

```
#Reformat the dataset so the rownames aren't their own column, however this is
#destructive
rownames(x) <- x[,1]
x <- x[,-1]
```

```
#Instead, let's read it in properly to begin with
x <- read.csv(url, row.names=1)
```


Now let's plot the data.

```
#Not an incredibly effective/readable plot
barplot(as.matrix(x), col=rainbow(17), beside=TRUE)
```



An exploratory plot that may be useful is pairs.

```
#You can visualize correlations between groups pairwise
mycols <-rainbow(nrow(x))
pairs(x, col=mycols, pch=16)
```



PCA to the rescue!


Here we will use the base R function for PCA, which is called `prcomp()`. This function wants the transpose of our data.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

##	PC1	PC2	PC3	PC4
## Standard deviation	324.1502	212.7478	73.87622	4.189e-14
## Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
## Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
plot(pca)
```




We want the score plot (a.k.a. PCA plot). Basically of PC1 vs PC2.

```
attributes(pca)
```

```
## $names
## [1] "sdev"      "rotation" "center"   "scale"    "x"
##
## $class
## [1] "prcomp"
```


We are after the `pca$x` component for this plot.

```
plot(pca$x[,1:2])
text(pca$x[,1:2], labels=colnames(x), col=c("orange", "pink", "blue", "green"))
```

We can also examine the PCA “loadings”, which tell us how much each food contributed to the principle component.

```
par(mar=c(10, 3, 0.35, 0))  
barplot(pca$rotation[,1], las=2)
```



One more PCA for today

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2  wt3  wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458  408  429 420  90  88  86  90  93
## gene2 219 200  204  210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4  783 792  829  856 760 849 856 835 885 894
## gene5  181 249  204  244 225 277 305 272 270 279
## gene6  460 502  491  491 493 612 594 577 618 638
```

```
nrow(rna.data)
```

```
## [1] 100
```

```
ncol(rna.data)
```

```
## [1] 10
```

```
colnames(rna.data)
```


```
## [1] "wt1" "wt2" "wt3" "wt4" "wt5" "ko1" "ko2" "ko3" "ko4" "ko5"
```

```
pca.rna = prcomp(t(rna.data), scale=TRUE)  
summary(pca.rna)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7  
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111  
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642  
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251  
##              PC8      PC9      PC10  
## Standard deviation  0.62065 0.60342 3.348e-15  
## Proportion of Variance 0.00385 0.00364 0.000e+00  
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```

```
plot(pca.rna)
```



```
plot(pca.rna$x[,1:2])  
text(pca.rna$x[,1:2], labels=colnames(rna.data))
```

