

Class09 Mini Project

Pierce Ford

10/27/2021

Unsupervised Learning Analysis of Human Breast Cancer Cells

```
#Read in the data file
fna.data <- "WisconsinCancer.csv"

#Convert data to data frame
wisc.df <- read.csv(fna.data, row.names=1)

#View the data to determine if the structure is as expected
str(wisc.df)
```

```
## 'data.frame':  569 obs. of  32 variables:
## $ diagnosis      : chr  "M" "M" "M" "M" ...
## $ radius_mean    : num  18 20.6 19.7 11.4 20.3 ...
## $ texture_mean    : num  10.4 17.8 21.2 20.4 14.3 ...
## $ perimeter_mean  : num  122.8 132.9 130 77.6 135.1 ...
## $ area_mean       : num  1001 1326 1203 386 1297 ...
## $ smoothness_mean : num  0.1184 0.0847 0.1096 0.1425 0.1003 ...
## $ compactness_mean : num  0.2776 0.0786 0.1599 0.2839 0.1328 ...
## $ concavity_mean  : num  0.3001 0.0869 0.1974 0.2414 0.198 ...
## $ concave.points_mean : num  0.1471 0.0702 0.1279 0.1052 0.1043 ...
## $ symmetry_mean   : num  0.242 0.181 0.207 0.26 0.181 ...
## $ fractal_dimension_mean : num  0.0787 0.0567 0.06 0.0974 0.0588 ...
## $ radius_se       : num  1.095 0.543 0.746 0.496 0.757 ...
## $ texture_se       : num  0.905 0.734 0.787 1.156 0.781 ...
## $ perimeter_se     : num  8.59 3.4 4.58 3.44 5.44 ...
## $ area_se          : num  153.4 74.1 94 27.2 94.4 ...
## $ smoothness_se    : num  0.0064 0.00522 0.00615 0.00911 0.01149 ...
## $ compactness_se   : num  0.049 0.0131 0.0401 0.0746 0.0246 ...
## $ concavity_se     : num  0.0537 0.0186 0.0383 0.0566 0.0569 ...
## $ concave.points_se : num  0.0159 0.0134 0.0206 0.0187 0.0188 ...
## $ symmetry_se      : num  0.03 0.0139 0.0225 0.0596 0.0176 ...
## $ fractal_dimension_se : num  0.00619 0.00353 0.00457 0.00921 0.00511 ...
## $ radius_worst     : num  25.4 25 23.6 14.9 22.5 ...
## $ texture_worst    : num  17.3 23.4 25.5 26.5 16.7 ...
## $ perimeter_worst  : num  184.6 158.8 152.5 98.9 152.2 ...
## $ area_worst       : num  2019 1956 1709 568 1575 ...
## $ smoothness_worst : num  0.162 0.124 0.144 0.21 0.137 ...
## $ compactness_worst : num  0.666 0.187 0.424 0.866 0.205 ...
```

```
## $ concavity_worst      : num  0.712 0.242 0.45 0.687 0.4 ...
## $ concave.points_worst : num  0.265 0.186 0.243 0.258 0.163 ...
## $ symmetry_worst       : num  0.46 0.275 0.361 0.664 0.236 ...
## $ fractal_dimension_worst: num  0.1189 0.089 0.0876 0.173 0.0768 ...
## $ X                     : logi  NA NA NA NA NA NA ...
```

```
head(wisc.df)
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean
## 842302      M      17.99      10.38      122.80      1001.0
## 842517      M      20.57      17.77      132.90      1326.0
## 84300903     M      19.69      21.25      130.00      1203.0
## 84348301      M      11.42      20.38      77.58      386.1
## 84358402      M      20.29      14.34      135.10     1297.0
## 843786      M      12.45      15.70      82.57      477.1
##      smoothness_mean compactness_mean concavity_mean concave.points_mean
## 842302      0.11840      0.27760      0.3001      0.14710
## 842517      0.08474      0.07864      0.0869      0.07017
## 84300903     0.10960      0.15990      0.1974      0.12790
## 84348301     0.14250      0.28390      0.2414      0.10520
## 84358402     0.10030      0.13280      0.1980      0.10430
## 843786      0.12780      0.17000      0.1578      0.08089
##      symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
## 842302      0.2419      0.07871      1.0950      0.9053      8.589
## 842517      0.1812      0.05667      0.5435      0.7339      3.398
## 84300903     0.2069      0.05999      0.7456      0.7869      4.585
## 84348301     0.2597      0.09744      0.4956      1.1560      3.445
## 84358402     0.1809      0.05883      0.7572      0.7813      5.438
## 843786      0.2087      0.07613      0.3345      0.8902      2.217
##      area_se smoothness_se compactness_se concavity_se concave.points_se
## 842302     153.40      0.006399      0.04904      0.05373      0.01587
## 842517      74.08      0.005225      0.01308      0.01860      0.01340
## 84300903     94.03      0.006150      0.04006      0.03832      0.02058
## 84348301     27.23      0.009110      0.07458      0.05661      0.01867
## 84358402     94.44      0.011490      0.02461      0.05688      0.01885
## 843786      27.19      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903     0.02250      0.004571      23.57      25.53
## 84348301     0.05963      0.009208      14.91      26.50
## 84358402     0.01756      0.005115      22.54      16.67
## 843786      0.02165      0.005082      15.47      23.75
##      perimeter_worst area_worst smoothness_worst compactness_worst
## 842302      184.60      2019.0      0.1622      0.6656
## 842517      158.80      1956.0      0.1238      0.1866
## 84300903     152.50      1709.0      0.1444      0.4245
## 84348301      98.87      567.7      0.2098      0.8663
## 84358402     152.20      1575.0      0.1374      0.2050
## 843786      103.40      741.6      0.1791      0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903     0.4504      0.2430      0.3613
```

```
## 84348301      0.6869      0.2575      0.6638
## 84358402      0.4000      0.1625      0.2364
## 843786       0.5355      0.1741      0.3985
##      fractal_dimension_worst X
## 842302      0.11890 NA
## 842517      0.08902 NA
## 84300903     0.08758 NA
## 84348301     0.17300 NA
## 84358402     0.07678 NA
## 843786      0.12440 NA
```

```
#Remove diagnosis column as that is essentially the "answer" our unsupervised
#learning will be looking for, preserve diagnosis as a factor vector for later
wisc.data <- wisc.df[,-1]
#Remove NA column
wisc.data <- wisc.data[,-length(colnames(wisc.data))]
head(wisc.data)
```

```
##      radius_mean texture_mean perimeter_mean area_mean smoothness_mean
## 842302      17.99      10.38      122.80      1001.0      0.11840
## 842517      20.57      17.77      132.90      1326.0      0.08474
## 84300903     19.69      21.25      130.00      1203.0      0.10960
## 84348301     11.42      20.38      77.58      386.1      0.14250
## 84358402     20.29      14.34      135.10      1297.0      0.10030
## 843786      12.45      15.70      82.57      477.1      0.12780
##      compactness_mean concavity_mean concave.points_mean symmetry_mean
## 842302      0.27760      0.3001      0.14710      0.2419
## 842517      0.07864      0.0869      0.07017      0.1812
## 84300903     0.15990      0.1974      0.12790      0.2069
## 84348301     0.28390      0.2414      0.10520      0.2597
## 84358402     0.13280      0.1980      0.10430      0.1809
## 843786      0.17000      0.1578      0.08089      0.2087
##      fractal_dimension_mean radius_se texture_se perimeter_se area_se
## 842302      0.07871      1.0950      0.9053      8.589 153.40
## 842517      0.05667      0.5435      0.7339      3.398 74.08
## 84300903     0.05999      0.7456      0.7869      4.585 94.03
## 84348301     0.09744      0.4956      1.1560      3.445 27.23
## 84358402     0.05883      0.7572      0.7813      5.438 94.44
## 843786      0.07613      0.3345      0.8902      2.217 27.19
##      smoothness_se compactness_se concavity_se concave.points_se
## 842302      0.006399      0.04904      0.05373      0.01587
## 842517      0.005225      0.01308      0.01860      0.01340
## 84300903     0.006150      0.04006      0.03832      0.02058
## 84348301     0.009110      0.07458      0.05661      0.01867
## 84358402     0.011490      0.02461      0.05688      0.01885
## 843786      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903     0.02250      0.004571      23.57      25.53
## 84348301     0.05963      0.009208      14.91      26.50
## 84358402     0.01756      0.005115      22.54      16.67
## 843786      0.02165      0.005082      15.47      23.75
##      perimeter_worst area_worst smoothness_worst compactness_worst
```

```
## 842302      184.60      2019.0      0.1622      0.6656
## 842517      158.80      1956.0      0.1238      0.1866
## 84300903     152.50      1709.0      0.1444      0.4245
## 84348301      98.87       567.7      0.2098      0.8663
## 84358402     152.20      1575.0      0.1374      0.2050
## 843786      103.40       741.6      0.1791      0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903     0.4504      0.2430      0.3613
## 84348301     0.6869      0.2575      0.6638
## 84358402     0.4000      0.1625      0.2364
## 843786      0.5355      0.1741      0.3985
##      fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903     0.08758
## 84348301     0.17300
## 84358402     0.07678
## 843786      0.12440
```

```
diagnosis <- as.vector(wisc.df$diagnosis)
diagnosis_factor <- factor(diagnosis)
```

Q1. How many observations are in this dataset?

```
#The number of observations is equal to the number of rows in the dataset
nrow(wisc.df)
```

```
## [1] 569
```

Q2. How many of the observations have a malignant diagnosis?

```
#Extract number of malignant samples using table
table(diagnosis_factor)["M"]
```

```
##      M
## 212
```

Q3. How many variables/features in the data are suffixed with `_mean`?

```
#Pull the columns that contain "_mean" and count them
mean_columns <- grep("_mean", colnames(wisc.data))
length(mean_columns)
```

```
## [1] 10
```

Principal Component Analysis

```
#Does the data need to be scaled? Check column means and standard deviations
colMeans(wisc.data)
```

```
##          radius_mean      texture_mean      perimeter_mean
##      1.412729e+01      1.928965e+01      9.196903e+01
##          area_mean      smoothness_mean      compactness_mean
##      6.548891e+02      9.636028e-02      1.043410e-01
##      concavity_mean      concave.points_mean      symmetry_mean
##      8.879932e-02      4.891915e-02      1.811619e-01
##      fractal_dimension_mean      radius_se      texture_se
##      6.279761e-02      4.051721e-01      1.216853e+00
##      perimeter_se      area_se      smoothness_se
##      2.866059e+00      4.033708e+01      7.040979e-03
##      compactness_se      concavity_se      concave.points_se
##      2.547814e-02      3.189372e-02      1.179614e-02
##      symmetry_se      fractal_dimension_se      radius_worst
##      2.054230e-02      3.794904e-03      1.626919e+01
##      texture_worst      perimeter_worst      area_worst
##      2.567722e+01      1.072612e+02      8.805831e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      1.323686e-01      2.542650e-01      2.721885e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      1.146062e-01      2.900756e-01      8.394582e-02
```

```
apply(wisc.data, 2, sd)
```

```
##          radius_mean      texture_mean      perimeter_mean
##      3.524049e+00      4.301036e+00      2.429898e+01
##          area_mean      smoothness_mean      compactness_mean
##      3.519141e+02      1.406413e-02      5.281276e-02
##      concavity_mean      concave.points_mean      symmetry_mean
##      7.971981e-02      3.880284e-02      2.741428e-02
##      fractal_dimension_mean      radius_se      texture_se
##      7.060363e-03      2.773127e-01      5.516484e-01
##      perimeter_se      area_se      smoothness_se
##      2.021855e+00      4.549101e+01      3.002518e-03
##      compactness_se      concavity_se      concave.points_se
##      1.790818e-02      3.018606e-02      6.170285e-03
##      symmetry_se      fractal_dimension_se      radius_worst
##      8.266372e-03      2.646071e-03      4.833242e+00
##      texture_worst      perimeter_worst      area_worst
##      6.146258e+00      3.360254e+01      5.693570e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      2.283243e-02      1.573365e-01      2.086243e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      6.573234e-02      6.186747e-02      1.806127e-02
```

The data has a wide range of standard deviations, so it should be scaled on a per column basis so the ones with higher variance don't automatically contribute more to the PCA.

```
#Run PCA and look at summary
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29     PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

As seen in the summary above, ~44% of the original variance is captured by PC1.

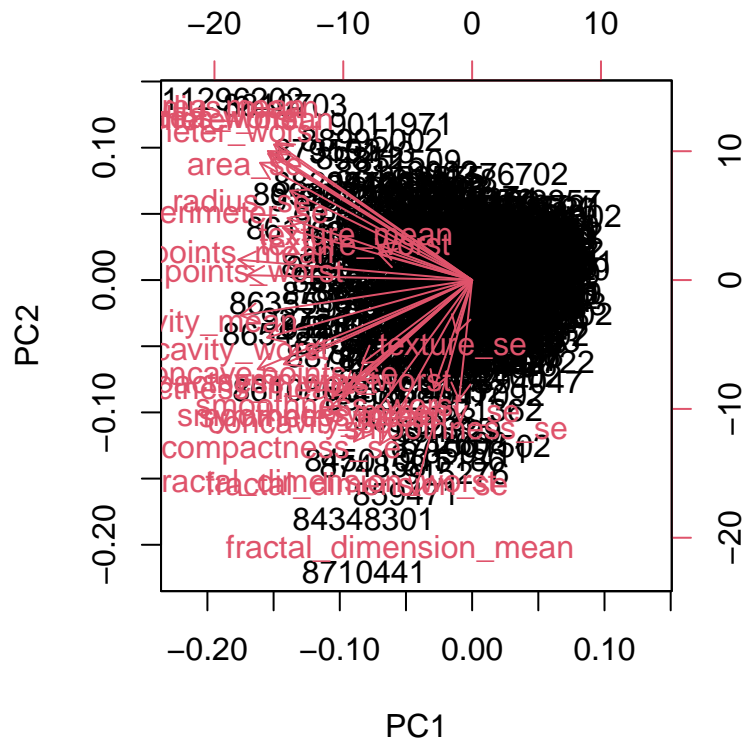
Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

At least 3 PCs are required to describe 70% of the variance (see cumulative proportion in summary).

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

At least 7 PCs are required to describe 90% of the variance (see cumulative proportion in summary).

```
#Let's plot the PCA!
biplot(wisc.pr)
```

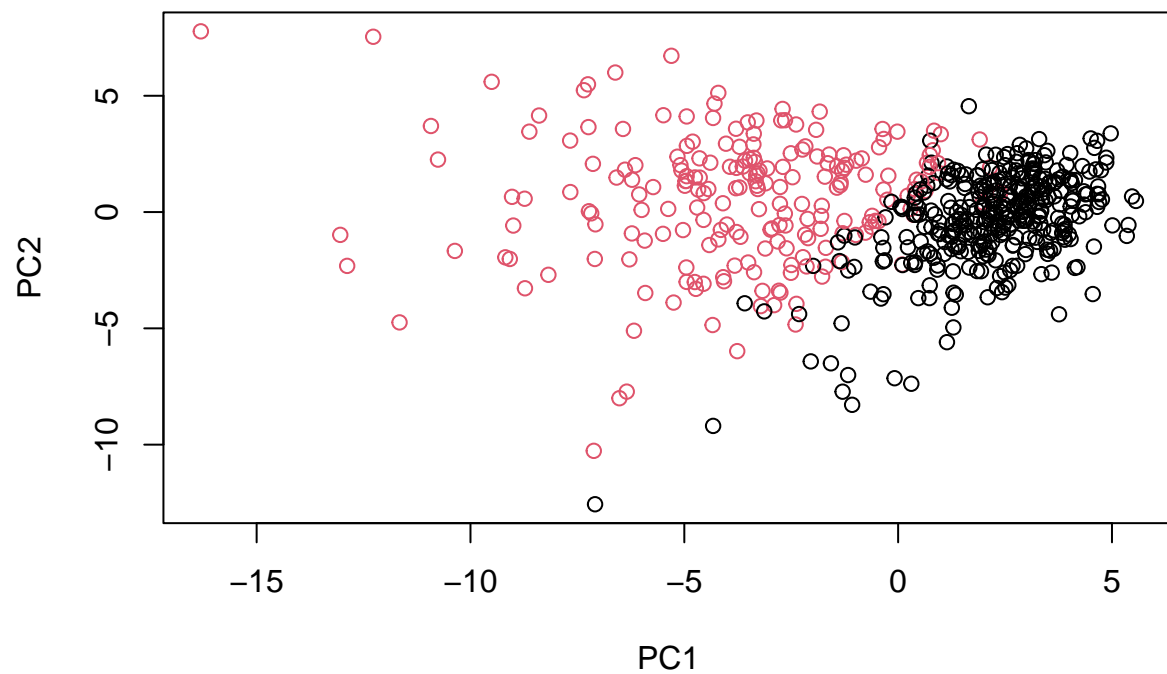


Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is impossible to read and cannot show us much of anything.

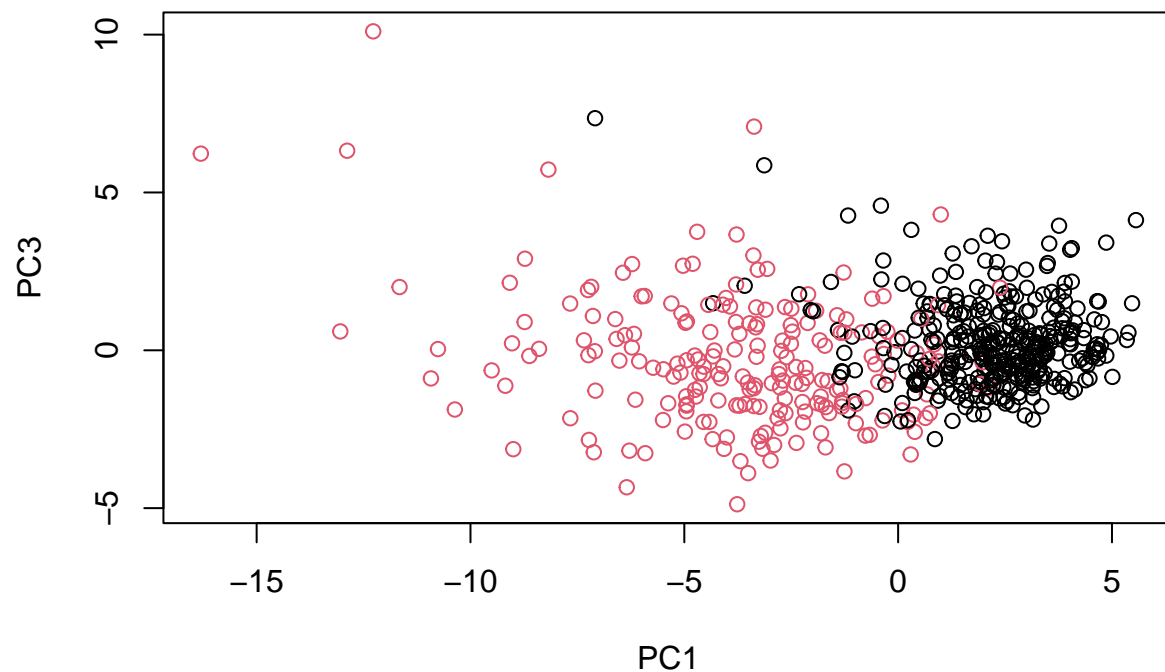
Instead, let's plot a scatter plot colored by diagnosis.

```
plot(wisc.pr$x[,1:2], col = diagnosis_factor,
     xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
# Repeat for components 1 and 3  
plot(wisc.pr$x[,c(1,3)], col = diagnosis_factor,  
     xlab = "PC1", ylab = "PC3")
```

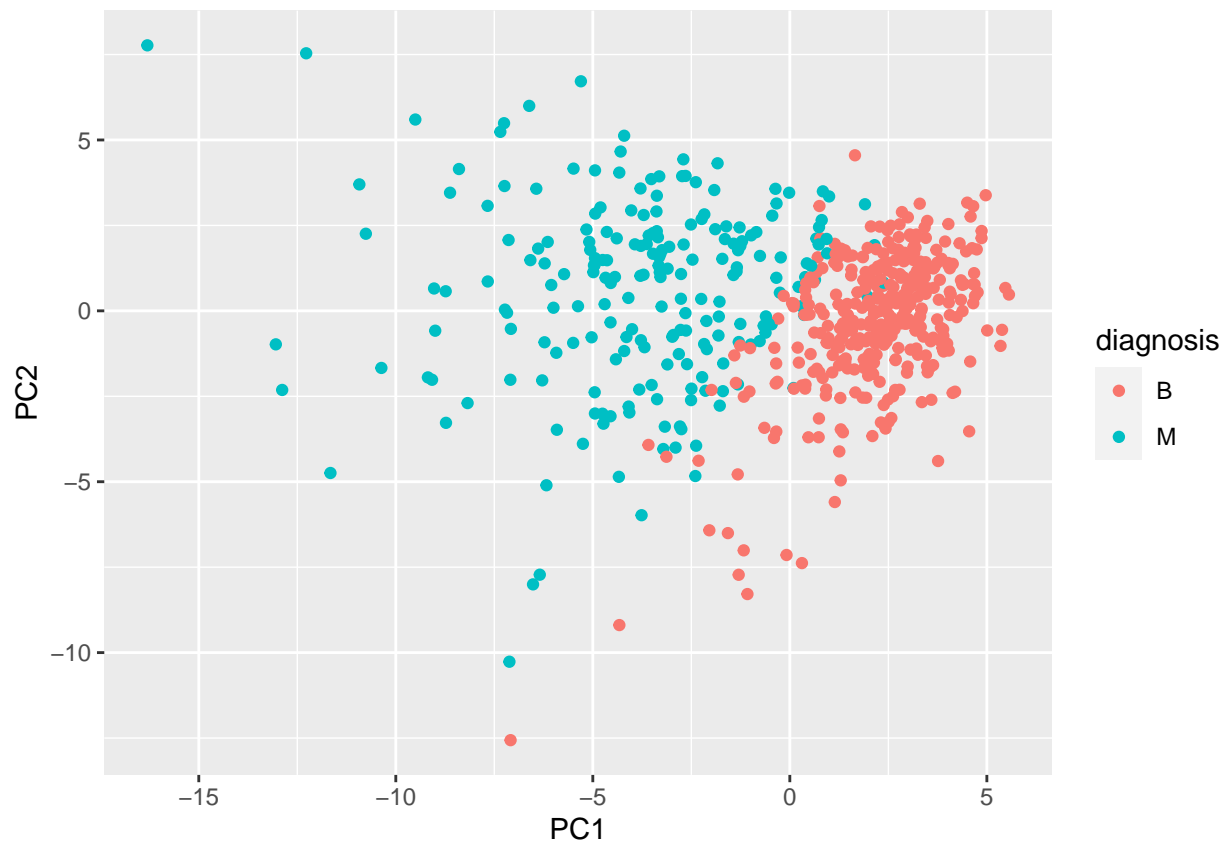
This plot has much poorer separation of the two clusters than the first, because PC3 captures less variance than PC2.

Making Nicer Plots With Ggplot

```
# Create a data.frame for ggplot
wisc.pr.df <- as.data.frame(wisc.pr$x)
#wisc.pr.df$diagnosis <- diagnosis_factor

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(wisc.pr.df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



What is the variance in each PC?

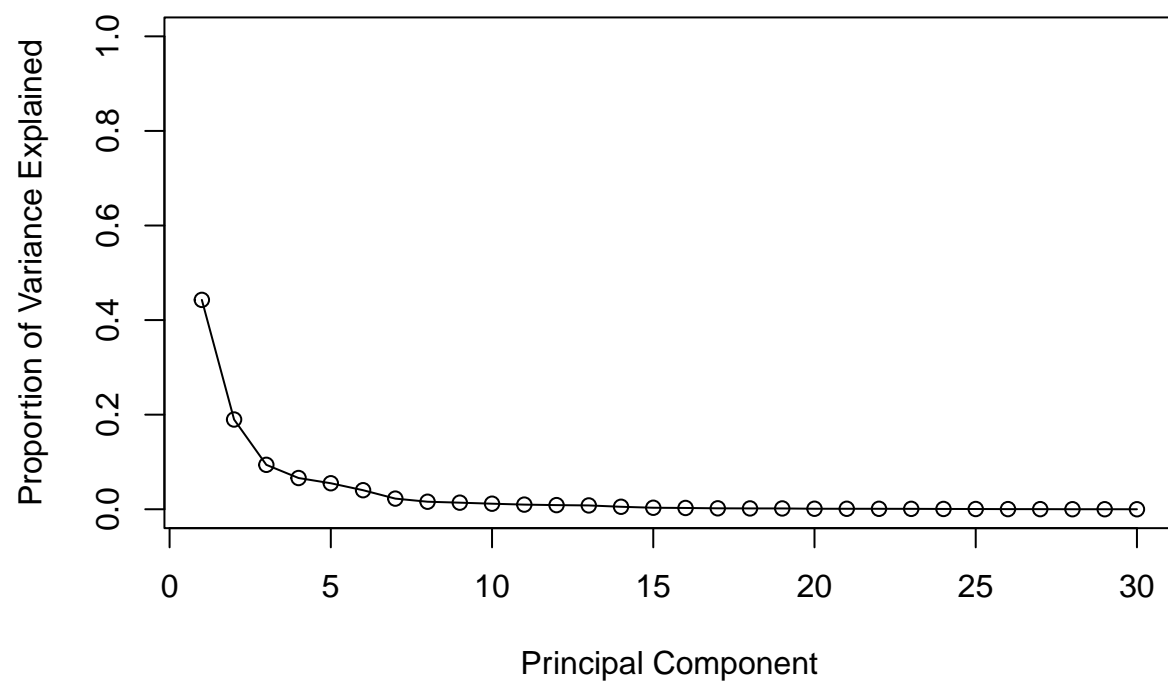
```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
pr.var
```

```
## [1] 1.328161e+01 5.691355e+00 2.817949e+00 1.980640e+00 1.648731e+00
## [6] 1.207357e+00 6.752201e-01 4.766171e-01 4.168948e-01 3.506935e-01
## [11] 2.939157e-01 2.611614e-01 2.413575e-01 1.570097e-01 9.413497e-02
## [16] 7.986280e-02 5.939904e-02 5.261878e-02 4.947759e-02 3.115940e-02
## [21] 2.997289e-02 2.743940e-02 2.434084e-02 1.805501e-02 1.548127e-02
## [26] 8.177640e-03 6.900464e-03 1.589338e-03 7.488031e-04 1.330448e-04
```

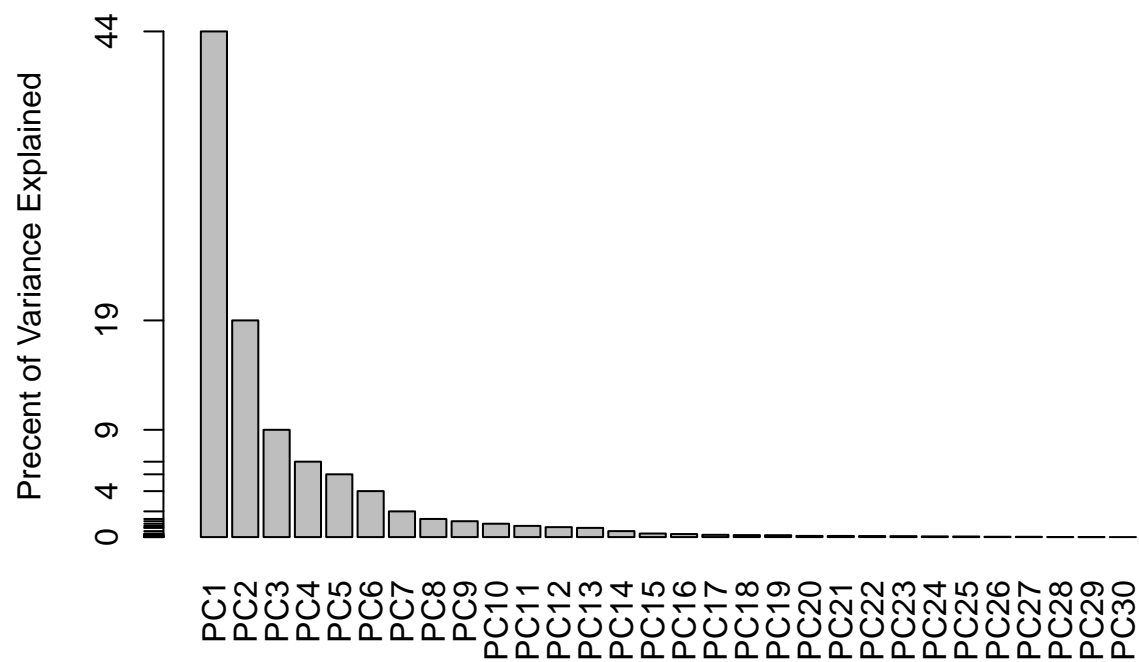
What is the proportion of variance explained by each PC? Let's look at some plots of this.

```
# Variance explained by each principal component: pve
pve <- pr.var/sum(pr.var)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



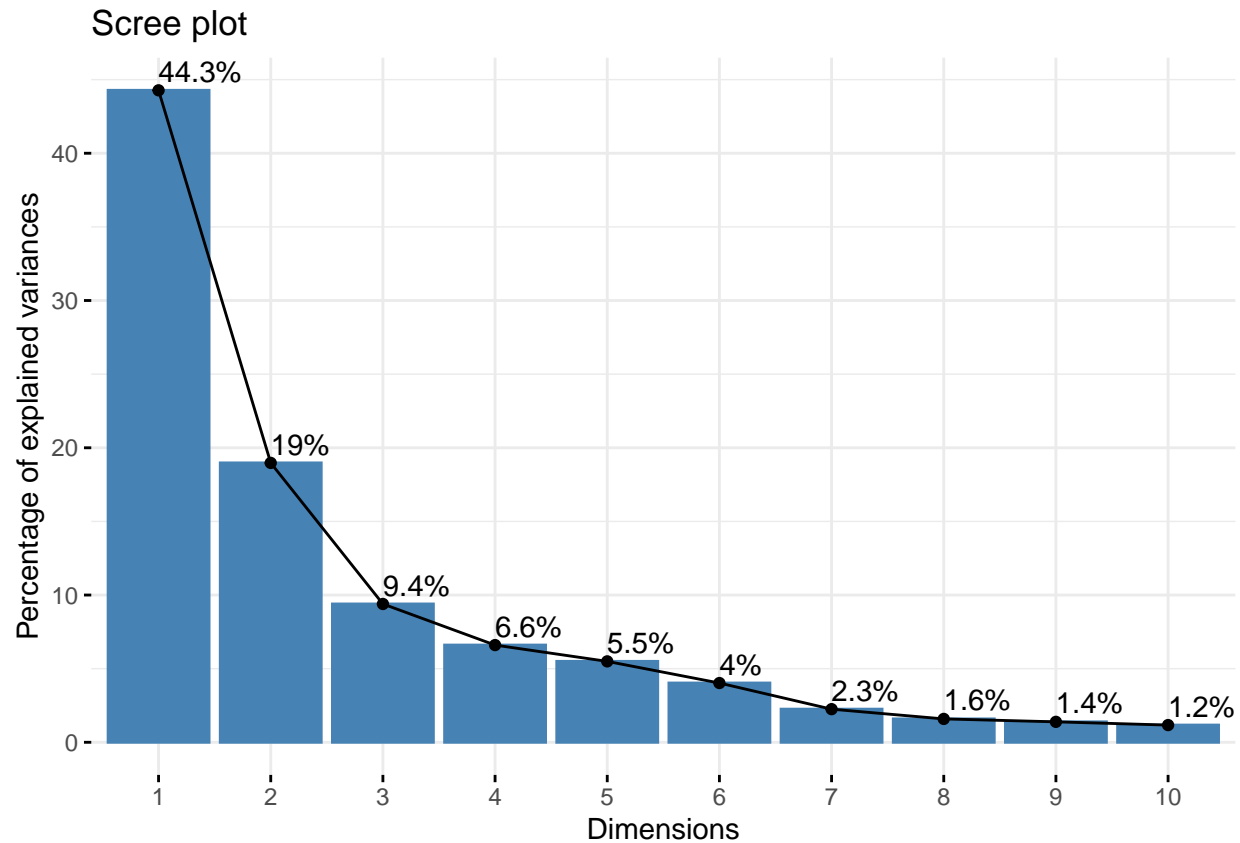
```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100)
```



```
#ggplot based graph
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
fviz_eig(wisc.pr, addlabels = TRUE)
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation[,1]["concave.points_mean"]
```

```
## concave.points_mean
## -0.2608538
```

#This number represents how much "concave.points_mean" contributes to PC1

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

```
#Sum each PC's variance until 80% is reached
sum <- 0
for (i in 1:length(pve)){
  sum <- sum + pve[i]
  if (sum >= 0.80){
    print(i)
    break
  }
}
```

```
## [1] 5
```

```
#Note this can be determined more easily by looking at the summary(wisc.pr)  
#results
```

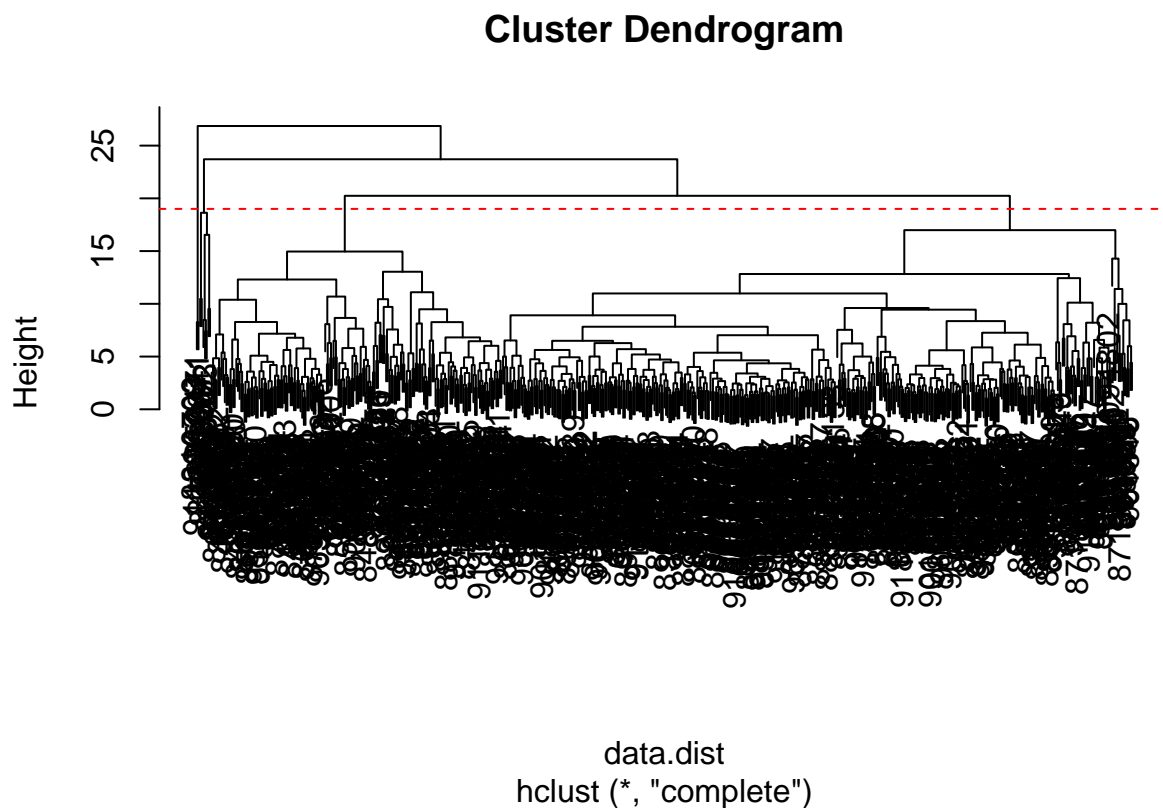
Hierarchical Clustering

```
# Scale the wisc.data data using the "scale()" function  
data.scaled <- scale(wisc.data)  
#Calculate distance (euclidean)  
data.dist <- dist(data.scaled)  
#Do the clustering  
wisc.hclust <- hclust(data.dist, method="complete")
```

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

Time to plot the cluster dendrogram.

```
plot(wisc.hclust)  
abline(h=19, col="red", lty=2)
```



A line height between 19 and 20 cuts the data into 4 clusters.

```
#We can also set the number of clusters
wisc.hclust.clusters <- cutree(wisc.hclust, 4)
#Compare clustes to diagnosis
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1  12 165
##              2   2   5
##              3 343  40
##              4   0   2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
#Let's try cutting into fewer clusters and see how well it matches diagnosis
#Three?
wisc.hclust.clusters <- cutree(wisc.hclust, 3)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1 355 205
##              2   2   5
##              3   0   2
```

```
#Two?
wisc.hclust.clusters <- cutree(wisc.hclust, 2)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1 357 210
##              2   0   2
```

```
#Both of these fail, with benign and malignant largely clustering together
#Let's try more
#Five?
wisc.hclust.clusters <- cutree(wisc.hclust, 5)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1  12 165
##              2   0   5
##              3 343  40
##              4   2   0
##              5   0   2
```

```
#Ten?
wisc.hclust.clusters <- cutree(wisc.hclust, 10)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1  12 86
##              2   0 59
##              3   0  3
##              4 331 39
##              5   0 20
##              6   2  0
##              7  12  0
##              8   0  2
##              9   0  2
##             10   0  1
```

#Even ten clusters can't separate fully into exclusively malignant and benign clusters, four clusters seems best

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
#Let's construct hclust objects using the various methods
wisc.hclust.complete <- hclust(data.dist, method="complete")
wisc.hclust.single <- hclust(data.dist, method="single")
wisc.hclust.average <- hclust(data.dist, method="average")
wisc.hclust.ward <- hclust(data.dist, method="ward.D2")

#Can any of the four give two good clusters?
wisc.hclust.complete.clusters_2 <- cutree(wisc.hclust.complete, 2)
table(wisc.hclust.complete.clusters_2, diagnosis)
```

```
##              diagnosis
## wisc.hclust.complete.clusters_2  B  M
##              1 357 210
##              2   0  2
```

```
wisc.hclust.single.clusters_2 <- cutree(wisc.hclust.single, 2)
table(wisc.hclust.single.clusters_2, diagnosis)
```

```
##              diagnosis
## wisc.hclust.single.clusters_2  B  M
##              1 357 210
##              2   0  2
```

```
wisc.hclust.average.clusters_2 <- cutree(wisc.hclust.average, 2)
table(wisc.hclust.average.clusters_2, diagnosis)
```

```
##              diagnosis
## wisc.hclust.average.clusters_2  B  M
##              1 357 209
##              2   0  3
```



```
wisc.hclust.ward.clusters_2 <- cutree(wisc.hclust.ward, 2)
table(wisc.hclust.ward.clusters_2, diagnosis)
```

```
##              diagnosis
## wisc.hclust.ward.clusters_2  B  M
##              1  20 164
##              2 337  48
```

```
#Ward clusters fairly well already with just two! How does each do with 4?
wisc.hclust.complete.clusters_4 <- cutree(wisc.hclust.complete, 4)
table(wisc.hclust.complete.clusters_4, diagnosis)
```

```
##              diagnosis
## wisc.hclust.complete.clusters_4  B  M
##              1  12 165
##              2   2   5
##              3 343  40
##              4   0   2
```

```
wisc.hclust.single.clusters_4 <- cutree(wisc.hclust.single, 4)
table(wisc.hclust.single.clusters_4, diagnosis)
```

```
##              diagnosis
## wisc.hclust.single.clusters_4  B  M
##              1 356 209
##              2   1   0
##              3   0   2
##              4   0   1
```

```
wisc.hclust.average.clusters_4 <- cutree(wisc.hclust.average, 4)
table(wisc.hclust.average.clusters_4, diagnosis)
```

```
##              diagnosis
## wisc.hclust.average.clusters_4  B  M
##              1 355 209
##              2   2   0
##              3   0   1
##              4   0   2
```

```
wisc.hclust.ward.clusters_4 <- cutree(wisc.hclust.ward, 4)
table(wisc.hclust.ward.clusters_4, diagnosis)
```

```
##              diagnosis
## wisc.hclust.ward.clusters_4  B  M
##              1   0 115
##              2   6  48
##              3 337  48
##              4  14   1
```

```
#Ward works about as well as complete does with 4 clusters, the other two are  
#nowhere near as good
```

The method “ward.D2” gives me my favorite results because it is the only one able to produce decent clusters when looking for 2 clusters, and the only one able to produce a pure cluster when looking for 4 clusters (cluster 1, all malignant).

K-means Clustering

```
#Run K-means on the data with 2 centers, 20 times  
wisc.km <- kmeans(data.scaled, centers= 2, nstart= 20)  
table(wisc.km$cluster, diagnosis_factor)
```

```
##      diagnosis_factor  
##           B      M  
##    1  14 175  
##    2 343  37
```

Q14. How well does k-means separate the two diagnoses? How does it compare to your hclust results?

The K-means method clusters the data slightly better (less “misdiagnosed” malignant samples in the benign cluster and vice versa as compared with the “ward.D2” method)

```
#How do the clusters relate?  
table(wisc.hclust.ward.clusters_2, wisc.km$cluster)
```

```
##  
## wisc.hclust.ward.clusters_2  1  2  
##                          1 173  11  
##                          2  16 369
```

```
table(wisc.hclust.ward.clusters_4, wisc.km$cluster)
```

```
##  
## wisc.hclust.ward.clusters_4  1  2  
##                          1 113  2  
##                          2  50  4  
##                          3  16 369  
##                          4  10  5
```

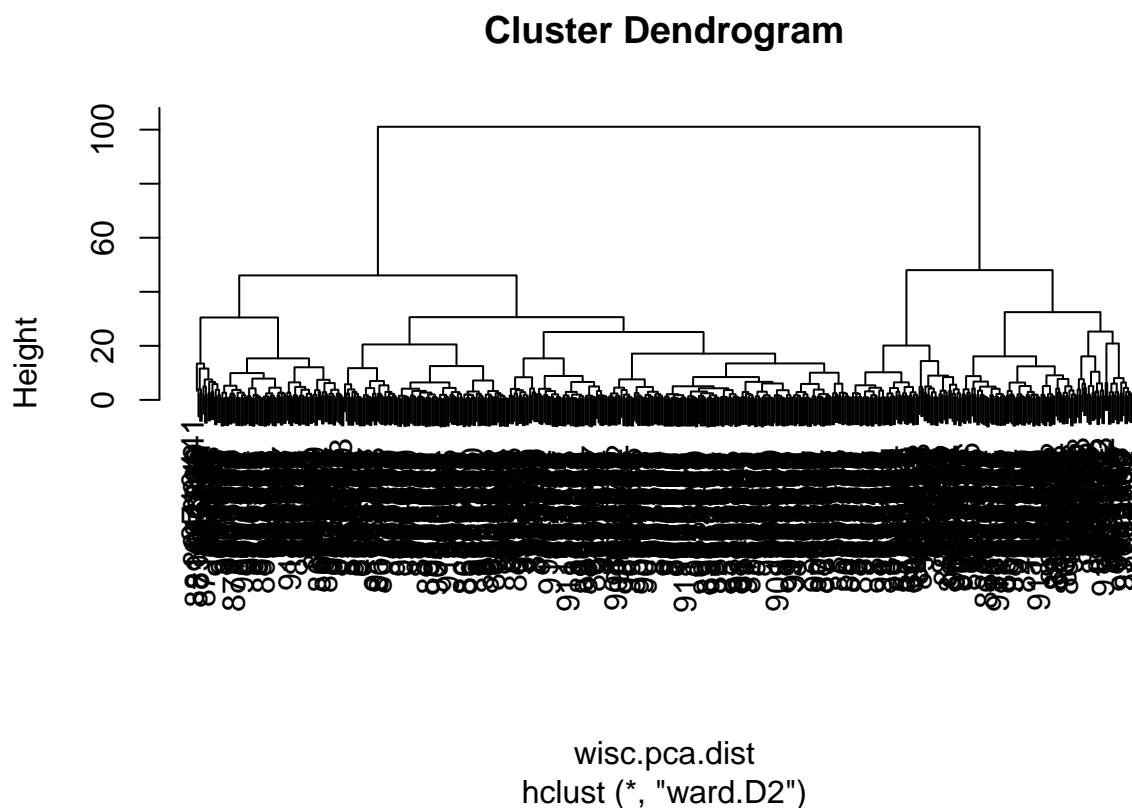
```
table(wisc.hclust.complete.clusters_4, wisc.km$cluster)
```

```
##  
## wisc.hclust.complete.clusters_4  1  2  
##                          1 160 17  
##                          2   7  0  
##                          3  20 363  
##                          4   2  0
```

```
#Each hclust clustering correlates fairly well with the k-means clustering
```

Combining Methods: Clustering with PCA results

```
#Run hclust on the PCA results with 4 PCs  
wisc.pca.dist <- dist(wisc.pr$x[,1:4])  
wisc.pr.hclust <- hclust(wisc.pca.dist, method="ward.D2")  
plot(wisc.pr.hclust)
```



```
#Group the results into 2 clusters  
grps <- cutree(wisc.pr.hclust, k=2)  
table(grps)
```

```
##      grps  
##      1   2  
## 171 398
```

```
#Compare grps to diagnosis  
table(grps, diagnosis_factor)
```

```
##      diagnosis_factor
```

```
## grps    B    M
##      1    6 165
##      2 351  47
```

#BONUS: how many PCs should we cluster on to minimize false negatives (min number of malignant samples in the benign cluster)

#Initialize variables

```
best_PCs <- 0
false_neg <- 400
```

#Test each number of PCs

```
for (i in 1:30){
```

#Cluster as before

```
temp.wisc.pca.dist <- dist(wisc.pr$x[,1:i])
temp.wisc.pr.hclust <- hclust(temp.wisc.pca.dist, method="ward.D2")
temp.grps <- cutree(temp.wisc.pr.hclust, k=2)
```

#Read in the false negatives

```
temp.false_neg <- min(table(temp.grps, diagnosis_factor)[1, "M"], table(temp.grps, diagnosis_factor)[
```

#Update if the nub=number of false negatives is fewer than the current min

```
if (temp.false_neg < false_neg){
  false_neg <- temp.false_neg
  best_PCs <- i
}
```

```
}
```

Q15. How well does the newly created model with two clusters (optimized for 10 PCs) separate out the two diagnoses?

#Output the number of PCs

```
print(best_PCs)
```

```
## [1] 10
```

#Produce clusters

```
wisc.pca.dist <- dist(wisc.pr$x[,1:best_PCs])
wisc.pr.hclust <- hclust(wisc.pca.dist, method="ward.D2")
grps <- cutree(wisc.pr.hclust, k=2)
```

#Output diagnosis table

```
table(grps, diagnosis_factor)
```

```
##      diagnosis_factor
## grps    B    M
##      1  39 205
##      2 318   7
```

The new model correlates very strongly with the proper diagnosis and minimizes the number of false negatives.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis
##      B      M
##  1  14  175
##  2 343   37
```

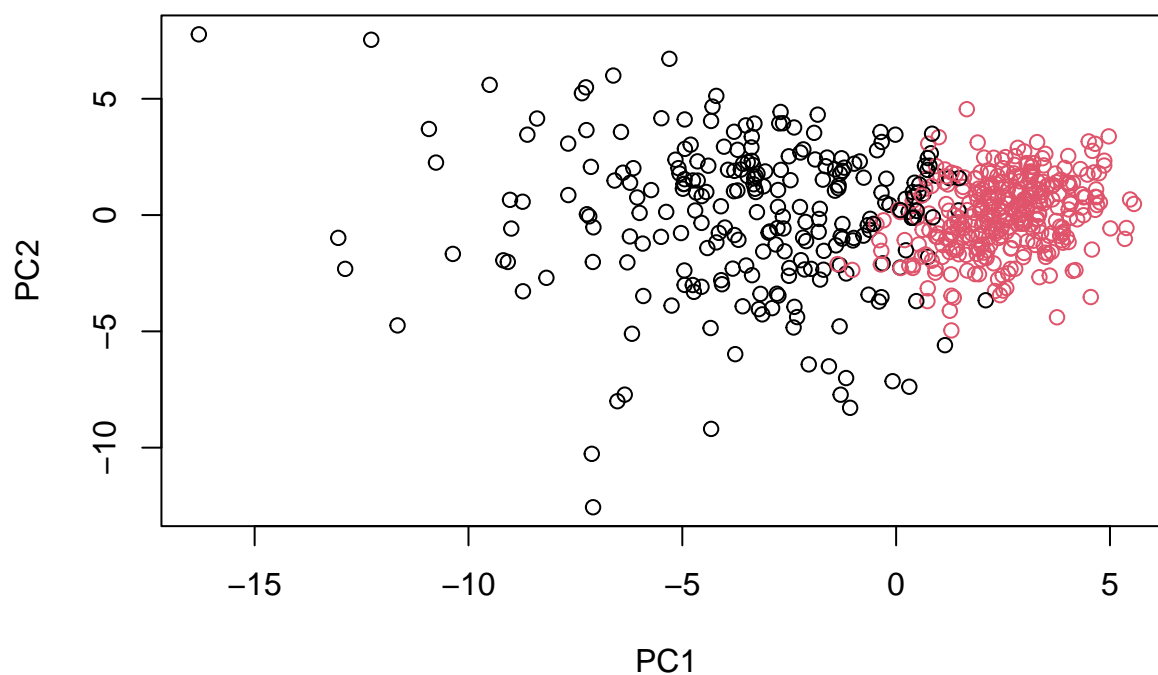
```
table(wisc.hclust.ward.clusters_2, diagnosis)
```

```
##              diagnosis
## wisc.hclust.ward.clusters_2  B      M
##              1  20  164
##              2 337   48
```

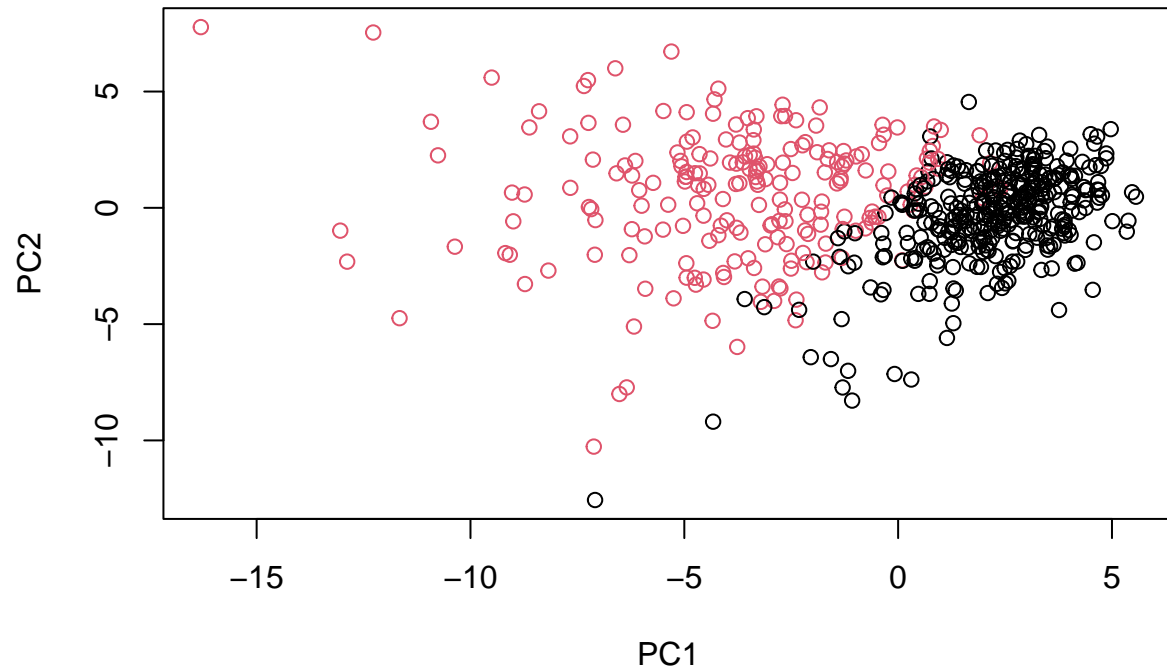
These two methods also correlate well, but the PCA clustering is more accurate and can be more easily fine-tuned to minimize false negatives.

For fun, let's plot the samples colored by the “best” grps and then by diagnosis.

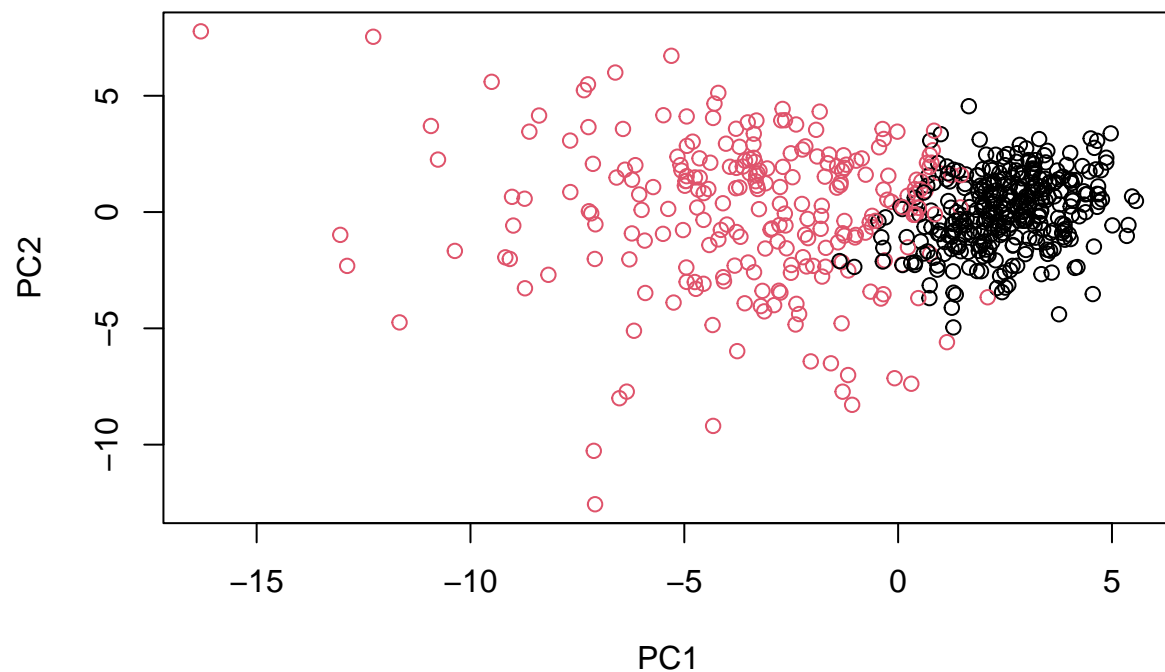
```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis_factor)
```



```
#Change colors to match  
g <- as.factor(grps)  
g <- relevel(g,2)  
plot(wisc.pr$x[,1:2], col=g)
```



Q17. Which of your analysis procedures resulted in a clustering model with the best specificity?
How about sensitivity?

```
#Calculate sensitivity for each model
```

```
#K-means sensitivity
```

```
Ksens <- 175/(175+37)
```

```
#Hclust sensitivity
```

```
Hsens <- 164/(164+48)
```

```
#PCA sensitivity
```

```
Psens <- 205/(205+7)
```

```
#What's best?
```

```
sens_vec <-c(Ksens=Ksens, Hsens=Hsens, Psens=Psens)
```

```
sens_vec
```

```
##      Ksens      Hsens      Psens
```

```
## 0.8254717 0.7735849 0.9669811
```

```
which.max(sens_vec)
```

```
## Psens
```

```
##      3
```

```
max(sens_vec)
```

```
## [1] 0.9669811
```

```
#Calculate specificity for each model
```

```
#K-means specificity
```

```
Kspec <- 343/(343+14)
```

```
#Hclust specificity
```

```
Hspec <- 337/(337+20)
```

```
#PCA specificity
```

```
Pspec <- 318/(318+39)
```

```
#What's best?
```

```
spec_vec <-c(Kspec=Kspec, Hspec=Hspec, Pspec=Pspect)
```

```
spec_vec
```

```
##      Kspec      Hspec      Pspec
```

```
## 0.9607843 0.9439776 0.8907563
```

```
which.max(spec_vec)
```

```
## Kspec
```

```
##      1
```

```
max(spec_vec)
```

```
## [1] 0.9607843
```

The PCA model is the most sensitive, but sacrifices specificity (k-means is best in terms of specificity, but I still believe the PCA is best for diagnostic purposes).

Can we now predict diagnoses on new data?

```
#Read in new samples
```

```
url <- "https://tinyurl.com/new-samples-CSV"
```

```
new <- read.csv(url)
```

```
npc <- predict(wisc.pr, newdata=new)
```

```
npc
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
```

```
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
```

```
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
```

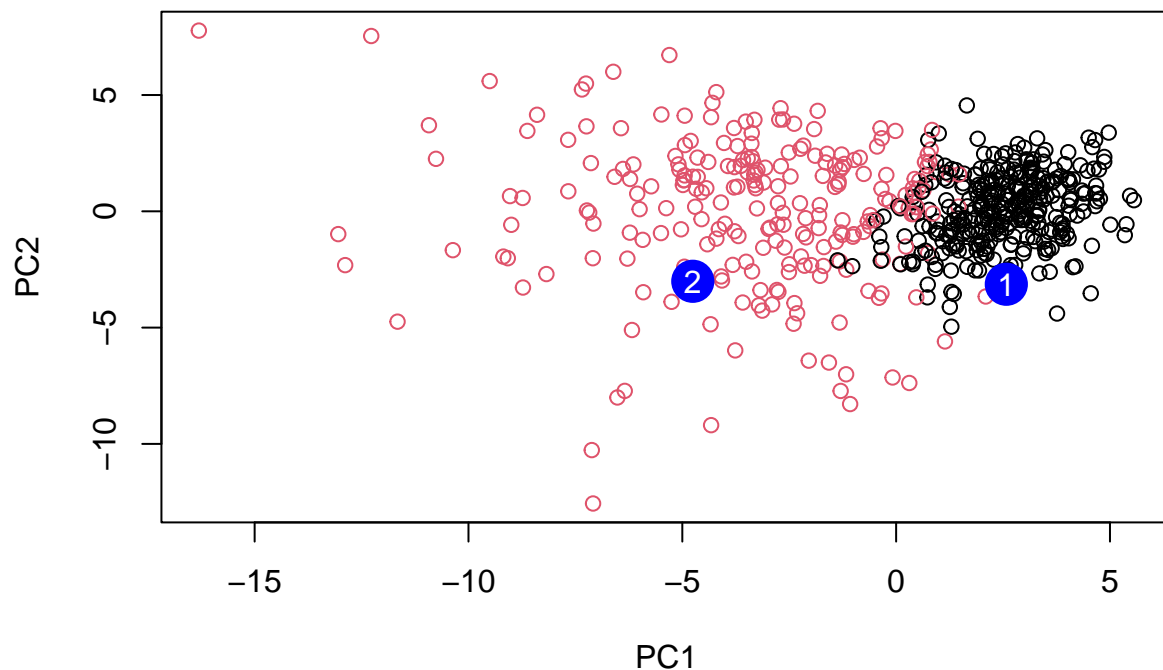
```
##          PC8          PC9          PC10          PC11          PC12          PC13          PC14
```

```
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
```



```
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##          PC15          PC16          PC17          PC18          PC19          PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
##          PC21          PC22          PC23          PC24          PC25          PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##          PC27          PC28          PC29          PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
#Plot the new samples
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

Patient 2 falls into the range of red points (majority malignant cluster), so they should be prioritized.